
MACHINE LEARNING PROJECT REPORT

DSBA

Contents

| | |
|--|-----------|
| 1.1 Read the dataset. Do the descriptive statistics and do the null value condition check. Write an inference on it. | 3 |
| 1.2 Perform Univariate and Bivariate Analysis. Do exploratory data analysis. Check for Outliers. ... | 5 |
| 1.3 Encode the data (having string values) for Modelling. Is Scaling necessary here or not? Data Split: Split the data into train and test (70:30). | 18 |
| 1.4 Apply Logistic Regression and LDA (linear discriminant analysis). | 22 |
| 1.5 Apply KNN Model and Naïve Bayes Model. Interpret the results..... | 29 |
| 1.6 Model Tuning, Bagging (Random Forest should be applied for Bagging) and Boosting. | 38 |
| 1.7 Performance Metrics: Check the performance of Predictions on Train and Test sets using Accuracy, Confusion Matrix, Plot ROC curve and get ROC_AUC score for each model. Final Model: Compare the models and write inference which model is best/optimized. | 46 |
| 1.8 Based on these predictions, what are the insights? | 64 |
| 2.1 Find the number of characters, words, and sentences for the mentioned documents. | 65 |
| 2.2) Remove all the stopwords from the three speeches. Show the word count before and after the removal of stopwords. Show a sample sentence after the removal of stopwords. | 65 |
| 2.3) Which word occurs the most number of times in his inaugural address for each president? Mention the top three words. (after removing the stopwords)..... | 66 |
| 2.4 Plot the word cloud of each of the speeches of the variable. (after removing the stopwords) . | 66 |

Problem 1:

You are hired by one of the leading news channels CNBE who wants to analyze recent elections. This survey was conducted on 1525 voters with 9 variables. You have to build a model, to predict which party a voter will vote for on the basis of the given information, to create an exit poll that will help in predicting overall win and seats covered by a particular party.

1.1 Read the dataset. Do the descriptive statistics and do the null value condition check. Write an inference on it.

- We have read the dataset given to us that contains the information of elections. The excel has 1525 rows and 9 columns.
- We have used the head () command to see the top 10 rows of the dataset shared as shown below:

| | Unnamed: 0 | vote | age | economic.cond.national | economic.cond.household | Blair | Hague | Europe | political.knowledge | gender |
|---|------------|--------|-----|------------------------|-------------------------|-------|-------|--------|---------------------|--------|
| 0 | 1 | Labour | 43 | 3 | 3 | 4 | 1 | 2 | 2 | female |
| 1 | 2 | Labour | 36 | 4 | 4 | 4 | 4 | 5 | 2 | male |
| 2 | 3 | Labour | 35 | 4 | 4 | 5 | 2 | 3 | 2 | male |
| 3 | 4 | Labour | 24 | 4 | 2 | 2 | 1 | 4 | 0 | female |
| 4 | 5 | Labour | 41 | 2 | 2 | 1 | 1 | 6 | 2 | male |

- We have used the tail () command to see the last 10 rows of the dataset shared as shown below:

| | Unnamed: 0 | vote | age | economic.cond.national | economic.cond.household | Blair | Hague | Europe | political.knowledge | gender |
|------|------------|--------------|-----|------------------------|-------------------------|-------|-------|--------|---------------------|--------|
| 1520 | 1521 | Conservative | 67 | 5 | 3 | 2 | 4 | 11 | 3 | male |
| 1521 | 1522 | Conservative | 73 | 2 | 2 | 4 | 4 | 8 | 2 | male |
| 1522 | 1523 | Labour | 37 | 3 | 3 | 5 | 4 | 2 | 2 | male |
| 1523 | 1524 | Conservative | 61 | 3 | 3 | 1 | 4 | 11 | 2 | male |
| 1524 | 1525 | Conservative | 74 | 2 | 3 | 2 | 4 | 11 | 0 | female |

- Dropped the unnamed column as it is of no use, please find the below dataset now:

| | vote | age | economic.cond.national | economic.cond.household | Blair | Hague | Europe | political.knowledge | gender |
|---|--------|-----|------------------------|-------------------------|-------|-------|--------|---------------------|--------|
| 0 | Labour | 43 | 3 | 3 | 4 | 1 | 2 | 2 | female |
| 1 | Labour | 36 | 4 | 4 | 4 | 4 | 5 | 2 | male |
| 2 | Labour | 35 | 4 | 4 | 5 | 2 | 3 | 2 | male |
| 3 | Labour | 24 | 4 | 2 | 2 | 1 | 4 | 0 | female |
| 4 | Labour | 41 | 2 | 2 | 1 | 1 | 6 | 2 | male |

- Information of the dataset

There is no null values as shown below:

```
>>> <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1525 entries, 0 to 1524
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Unnamed: 0          1525 non-null   int64
1   vote                1525 non-null   object
2   age                 1525 non-null   int64
3   economic.cond.national 1525 non-null   int64
4   economic.cond.household 1525 non-null   int64
5   Blair               1525 non-null   int64
6   Hague               1525 non-null   int64
7   Europe              1525 non-null   int64
8   political.knowledge  1525 non-null   int64
9   gender              1525 non-null   object
dtypes: int64(8), object(2)
memory usage: 119.3+ KB
```

- Shape of the dataset i.e. how many rows and columns are there in the shared dataset.

So we get 1525 rows and 9 columns

```
(1525, 9)
```

- Checked the null values present in dataset.

```
>>> vote                0
age                    0
economic.cond.national 0
economic.cond.household 0
Blair                  0
Hague                  0
Europe                 0
political.knowledge     0
gender                 0
dtype: int64
```

- Description of the data

| | count | mean | std | min | 25% | 50% | 75% | max |
|-------------------------|--------|-----------|-----------|------|------|------|------|------|
| age | 1525.0 | 54.182295 | 15.711209 | 24.0 | 41.0 | 53.0 | 67.0 | 93.0 |
| economic.cond.national | 1525.0 | 3.245902 | 0.880969 | 1.0 | 3.0 | 3.0 | 4.0 | 5.0 |
| economic.cond.household | 1525.0 | 3.140328 | 0.929951 | 1.0 | 3.0 | 3.0 | 4.0 | 5.0 |
| Blair | 1525.0 | 3.334426 | 1.174824 | 1.0 | 2.0 | 4.0 | 4.0 | 5.0 |
| Hague | 1525.0 | 2.746885 | 1.230703 | 1.0 | 2.0 | 2.0 | 4.0 | 5.0 |
| Europe | 1525.0 | 6.728525 | 3.297538 | 1.0 | 4.0 | 6.0 | 10.0 | 11.0 |
| political.knowledge | 1525.0 | 1.542295 | 1.083315 | 0.0 | 0.0 | 2.0 | 2.0 | 3.0 |

- Datatypes of the dataset given

```

vote                int8
age                 int64
economic.cond.national  int64
economic.cond.household int64
Blair               int64
Hague              int64
Europe             int64
political.knowledge int64
gender             int8
dtype: object

```

It could be noticed that there isn't much outlier considering the spread on the either side of median for continuous variables and that the target variable vote has only 2 distinct values indicating only two classes from classification perspective with Labour holding higher frequency. Also it could be noticed that gender has 2 unique values with female holding higher frequency of rows.

Refer the below exact distribution of the categorical variable:

```

VOTE : 2
Conservative    462
Labour          1063
Name: vote, dtype: int64

GENDER : 2
male        713
female      812
Name: gender, dtype: int64

```

Inferences:

1. The count of the dataset is 1525 row and 9 column
2. There is 8 duplicate we found in the dataset.
3. We have removed one column named un-named which consist of indices and that will be no use.
4. The mean age is found to be 54 and min age found is 24 as per analysis.
5. There is no null values present in dataset.
6. It could be noticed that there isn't much outlier considering the spread on the either side of median for continuous variables and that the target variable vote has only 2 distinct values

1.2 Perform Univariate and Bivariate Analysis. Do exploratory data analysis. Check for Outliers.

- Checked the null values present in dataset.

```

⇒ vote      0
   age      0
   economic.cond.national  0
   economic.cond.household  0
   Blair     0
   Hague     0
   Europe    0
   political.knowledge     0
   gender    0
   dtype: int64

```

As checked, there is no null values present in dataset.

- Description of the data

| | count | mean | std | min | 25% | 50% | 75% | max |
|-------------------------|--------|-----------|-----------|------|------|------|------|------|
| age | 1525.0 | 54.182295 | 15.711209 | 24.0 | 41.0 | 53.0 | 67.0 | 93.0 |
| economic.cond.national | 1525.0 | 3.245902 | 0.880969 | 1.0 | 3.0 | 3.0 | 4.0 | 5.0 |
| economic.cond.household | 1525.0 | 3.140328 | 0.929951 | 1.0 | 3.0 | 3.0 | 4.0 | 5.0 |
| Blair | 1525.0 | 3.334426 | 1.174824 | 1.0 | 2.0 | 4.0 | 4.0 | 5.0 |
| Hague | 1525.0 | 2.746885 | 1.230703 | 1.0 | 2.0 | 2.0 | 4.0 | 5.0 |
| Europe | 1525.0 | 6.728525 | 3.297538 | 1.0 | 4.0 | 6.0 | 10.0 | 11.0 |
| political.knowledge | 1525.0 | 1.542295 | 1.083315 | 0.0 | 0.0 | 2.0 | 2.0 | 3.0 |

- Datatypes of the dataset given

```

vote      int8
age      int64
economic.cond.national  int64
economic.cond.household  int64
Blair     int64
Hague     int64
Europe    int64
political.knowledge     int64
gender    int8
dtype: object

```


- Shape of the dataset i.e. how many rows and columns are there in the shared dataset.

So we get 1525 rows and 9 columns

```
(1525, 9)
```


UNIVARIATE ANALYSIS:













BOXPLOT:




We found that there are outliers present in 2 columns i.e. economic.cond.national and economic.cond.household .

Another display pattern of boxplot:






BI-VARIATE ANALYSIS:




Inference:

- The plot above shows that lower age group tend to vote for Labour while Labour also retains better vote share for higher age group comparatively.




Inference:

- There is no noticeable difference seen between these two parameters that is vote and economic.cond.household




Inference:

- The higher economic condition at national level condition of Labour party is higher as compares to Conservative.




Inference:

- Blair has received better higher scale assessment ratings among Labour voters compared to Conservative.

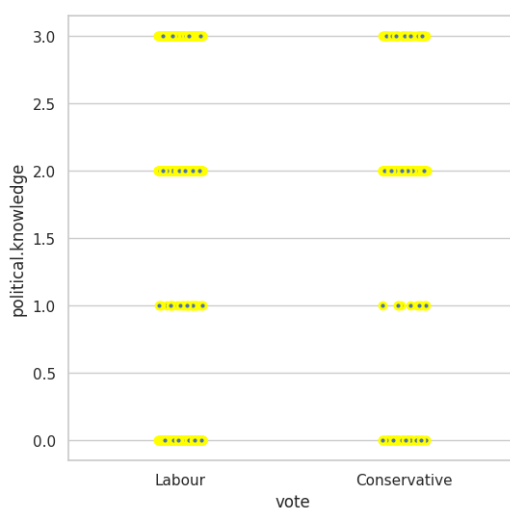


Inference:

- The above plots shows that Hague received higher scale ratings from Conservative voters while increased frequency in lower scale ratings from Labour




- The above plot shows that the Conservative party has a smaller number of voters with Eurosceptic sentiment.



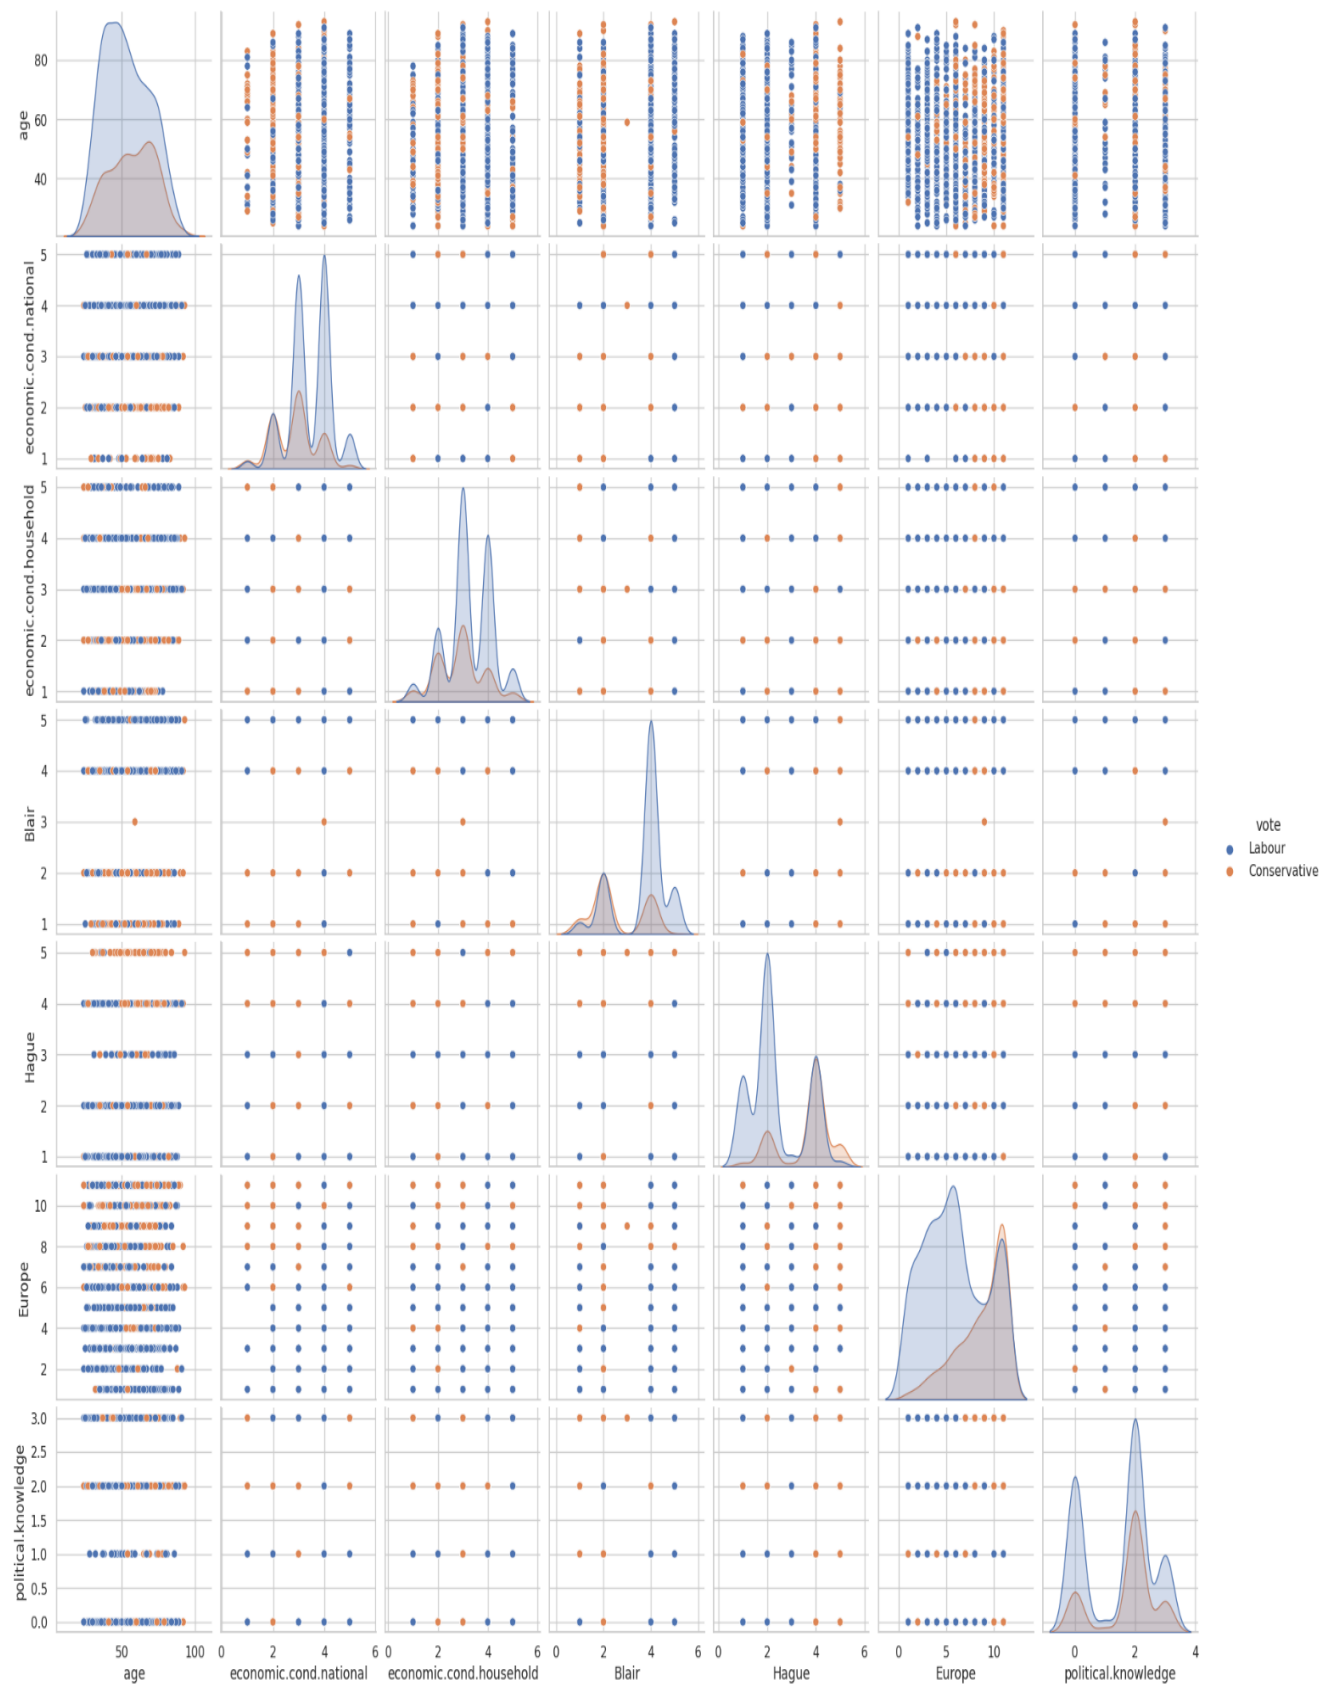
- Lesser number of voters for Conservative falls under lower scale knowledge regarding party position in European integrations.

MULTI-VARIATE ANALYSIS:

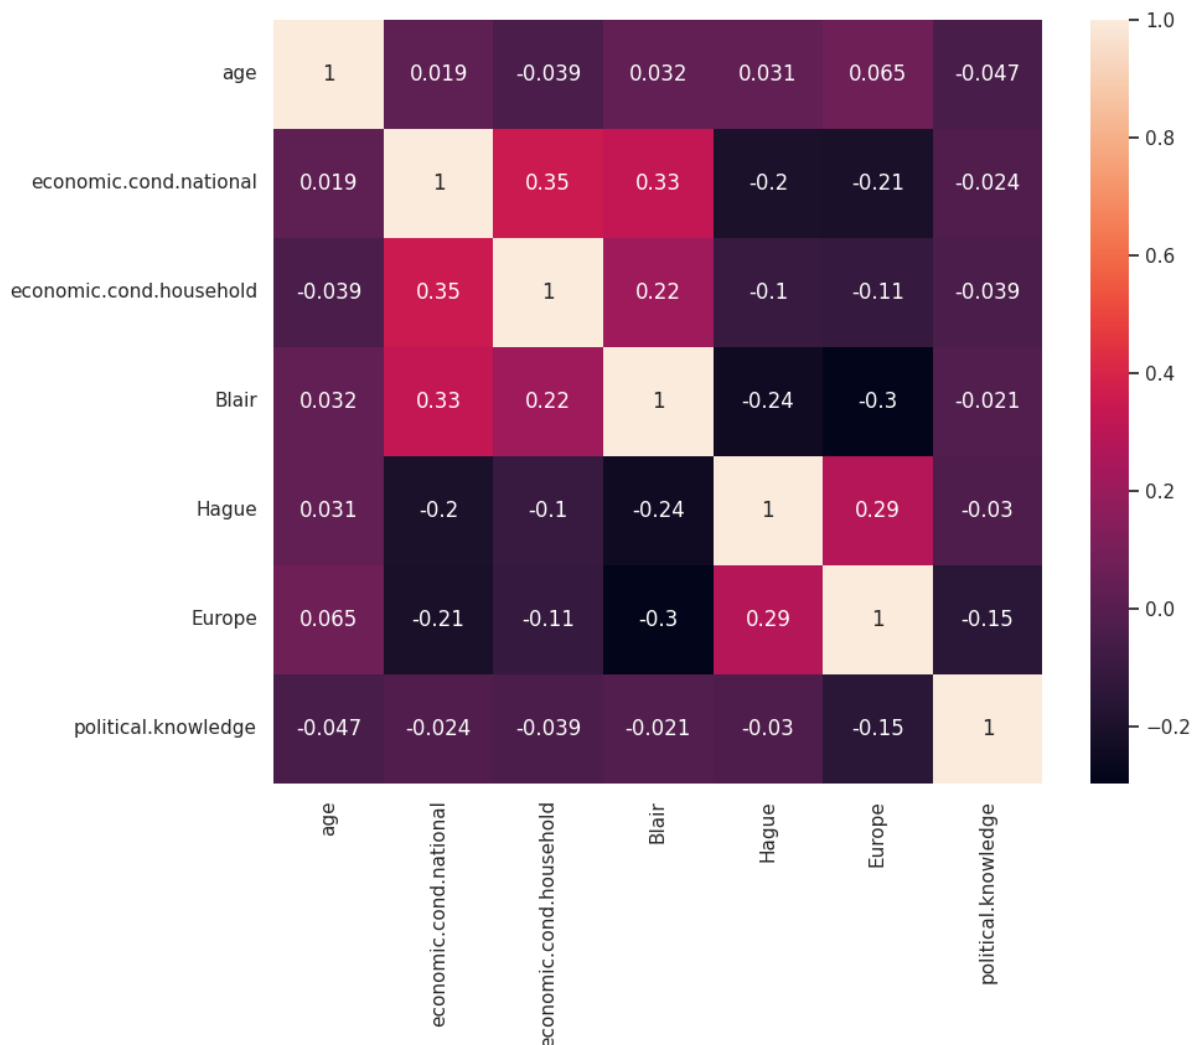
- Pair-plot for unscaled dataset



- Pair-plot using vote (labour or conservative) party as a hue.




- Heat Map




- There is no correlation between any variable.
- We have noticed that at lower age group is are less in frequency within lower scale range for national level economic condition. Also folks with highest political knowledge tends to dip with highest age group

We have treated the outliers present in the two columns as shown above . We have treated outliers by IQR Z score method . After applying the same we have removed the outliers

Before treating outliers:




More clear visualization:




Please note that both the boxplot are denoting the same thing. The only difference is visualization.

After treating outliers:



More clear visualization:



1.3 Encode the data (having string values) for Modelling. Is Scaling necessary here or not? Data Split: Split the data into train and test (70:30).

The target variable vote has only 2 distinct values indicating only two classes from classification perspective with Labour holding higher frequency. Also it could be noticed that gender has 2 unique values with female holding higher frequency of rows.

Refer the below exact distribution of the categorical variable:

```
VOTE : 2
Conservative    462
Labour          1063
Name: vote, dtype: int64

GENDER : 2
male        713
female      812
Name: gender, dtype: int64
```

We need to convert the categorical variables that is the gender column and Vote column to Numerical data to proceed further with the dataset.

We have given the below code for respective:

- In VOTE

Conservative -----0

Labour -----1

In gender column we have applied one hot encoding has been applied to gender to create binary variable depicting the binary status for each of the gender (Male/Female) and to optimize the first one has been dropped as the 0 or 1 in the other is sufficient to represent both.

| | vote | age | economic_cond_national | economic_cond_household | Blair | Hague | Europe | political_knowledge | gender_male |
|---|------|-----|------------------------|-------------------------|-------|-------|--------|---------------------|-------------|
| 0 | 1 | 43 | 3 | 3 | 4 | 1 | 2 | 2 | 0 |
| 1 | 1 | 36 | 4 | 4 | 4 | 4 | 5 | 2 | 1 |
| 2 | 1 | 35 | 4 | 4 | 5 | 2 | 3 | 2 | 1 |
| 3 | 1 | 24 | 4 | 2 | 2 | 1 | 4 | 0 | 0 |
| 4 | 1 | 41 | 2 | 2 | 1 | 1 | 6 | 2 | 1 |

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1517 entries, 0 to 1524
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   vote                                1517 non-null   int64
1   age                                1517 non-null   int64
2   economic_cond_national              1517 non-null   int64
3   economic_cond_household             1517 non-null   int64
4   Blair                              1517 non-null   int64
5   Hague                              1517 non-null   int64
6   Europe                             1517 non-null   int64
7   political_knowledge                 1517 non-null   int64
8   gender_male                         1517 non-null   uint8
dtypes: int64(8), uint8(1)
memory usage: 108.1 KB

```

| | vote | age | economic_cond_national | economic_cond_household | Blair | Hague | Europe | political_knowledge | gender_male |
|-------|-------------|-------------|------------------------|-------------------------|-------------|-------------|-------------|---------------------|-------------|
| count | 1517.000000 | 1517.000000 | 1517.000000 | 1517.000000 | 1517.000000 | 1517.000000 | 1517.000000 | 1517.000000 | 1517.000000 |
| mean | 0.696770 | 0.438279 | 0.561305 | 0.534443 | 0.583883 | 0.437376 | 0.574028 | 0.513514 | 0.467370 |
| std | 0.459805 | 0.227561 | 0.220448 | 0.232767 | 0.293693 | 0.308120 | 0.329904 | 0.361472 | 0.499099 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.000000 | 0.246377 | 0.500000 | 0.500000 | 0.250000 | 0.250000 | 0.300000 | 0.000000 | 0.000000 |
| 50% | 1.000000 | 0.420290 | 0.500000 | 0.500000 | 0.750000 | 0.250000 | 0.500000 | 0.666667 | 0.000000 |
| 75% | 1.000000 | 0.623188 | 0.750000 | 0.750000 | 0.750000 | 0.750000 | 0.900000 | 0.666667 | 1.000000 |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |

- **Data before Scaling:**

| | vote | age | economic_cond_national | economic_cond_household | Blair | Hague | Europe | political_knowledge | gender_male |
|---|------|-----|------------------------|-------------------------|-------|-------|--------|---------------------|-------------|
| 0 | 1 | 43 | 3 | 3 | 4 | 1 | 2 | 2 | 0 |
| 1 | 1 | 36 | 4 | 4 | 4 | 4 | 5 | 2 | 1 |
| 2 | 1 | 35 | 4 | 4 | 5 | 2 | 3 | 2 | 1 |
| 3 | 1 | 24 | 4 | 2 | 2 | 1 | 4 | 0 | 0 |
| 4 | 1 | 41 | 2 | 2 | 1 | 1 | 6 | 2 | 1 |

Scaling must be done as you can see there is major difference between the column of age along with other columns. However, scaling is a must step in

some model like KNN as it works on distance computation (Euclidean). However, for logistic regression as it goes by linear equation scaling must be done.


For LDA, scaling standardizes the coefficients of independent variables which helps in clear separation of classes as comparison of coefficients happens on standardized data.

As we know that Naive bayes is unaffected by scaling. Going by the fact that independent variables have different units, scaling becomes necessary to remove the units the variables are associated with so that the linear equation can be formed on the independent variables post standardization of them.




Z scoring based scaling of data would change the coefficient, neutralize/remove the intercept while the accuracy score remains the same before and after. MSE would get scaled too.

So we have used MIN MAX Scaler for scaling the data so that it will wouldn't affect the same.

- **Dataset after scaling:**



| | vote | age | economic_cond_national | economic_cond_household | Blair | Hague | Europe | political_knowledge | gender_male |
|------|------|----------|------------------------|-------------------------|-------|-------|--------|---------------------|-------------|
| 0 | 1.0 | 0.275362 | 0.50 | 0.50 | 0.75 | 0.00 | 0.1 | 0.666667 | 0.0 |
| 1 | 1.0 | 0.173913 | 0.75 | 0.75 | 0.75 | 0.75 | 0.4 | 0.666667 | 1.0 |
| 2 | 1.0 | 0.159420 | 0.75 | 0.75 | 1.00 | 0.25 | 0.2 | 0.666667 | 1.0 |
| 3 | 1.0 | 0.000000 | 0.75 | 0.25 | 0.25 | 0.00 | 0.3 | 0.000000 | 0.0 |
| 4 | 1.0 | 0.246377 | 0.25 | 0.25 | 0.00 | 0.00 | 0.5 | 0.666667 | 1.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1512 | 0.0 | 0.623188 | 1.00 | 0.50 | 0.25 | 0.75 | 1.0 | 1.000000 | 1.0 |
| 1513 | 0.0 | 0.710145 | 0.25 | 0.25 | 0.75 | 0.75 | 0.7 | 0.666667 | 1.0 |
| 1514 | 1.0 | 0.188406 | 0.50 | 0.50 | 1.00 | 0.75 | 0.1 | 0.666667 | 1.0 |
| 1515 | 0.0 | 0.536232 | 0.50 | 0.50 | 0.00 | 0.75 | 1.0 | 0.666667 | 1.0 |
| 1516 | 0.0 | 0.724638 | 0.25 | 0.50 | 0.25 | 0.75 | 1.0 | 0.000000 | 0.0 |





Data Split: Split the data into train and test (70:30)

The data has been split into 70 :30 train and test data for further analysis.

```
Number of rows and columns of the training set for the independent variables: (1067, 8)
Number of rows and columns of the training set for the dependent variable: (1067,)
Number of rows and columns of the test set for the independent variables: (458, 8)
Number of rows and columns of the test set for the dependent variable: (458,)
```

5 point summary for train data

| | count | mean | std | min | 25% | 50% | 75% | max |  |
|-------------------------|--------|----------|----------|-----|---------|----------|----------|-----|---|
| const | 1061.0 | 1.000000 | 0.000000 | 1.0 | 1.00000 | 1.000000 | 1.000000 | 1.0 |  |
| age | 1061.0 | 0.436763 | 0.225165 | 0.0 | 0.26087 | 0.420290 | 0.623188 | 1.0 | |
| economic_cond_national | 1061.0 | 0.561970 | 0.214428 | 0.0 | 0.50000 | 0.500000 | 0.750000 | 1.0 | |
| economic_cond_household | 1061.0 | 0.538407 | 0.234666 | 0.0 | 0.50000 | 0.500000 | 0.750000 | 1.0 | |
| Blair | 1061.0 | 0.589303 | 0.292155 | 0.0 | 0.25000 | 0.750000 | 0.750000 | 1.0 | |
| Hague | 1061.0 | 0.432846 | 0.308340 | 0.0 | 0.25000 | 0.250000 | 0.750000 | 1.0 | |
| Europe | 1061.0 | 0.565881 | 0.326047 | 0.0 | 0.30000 | 0.500000 | 0.900000 | 1.0 | |
| political_knowledge | 1061.0 | 0.504555 | 0.358639 | 0.0 | 0.00000 | 0.666667 | 0.666667 | 1.0 | |
| gender_male | 1061.0 | 0.467484 | 0.499177 | 0.0 | 0.00000 | 0.000000 | 1.000000 | 1.0 | |

5-point summary for test data

| | count | mean | std | min | 25% | 50% | 75% | max |  |
|-------------------------|-------|----------|----------|-----|----------|----------|----------|----------|---|
| const | 456.0 | 1.000000 | 0.000000 | 1.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |  |
| age | 456.0 | 0.441807 | 0.233256 | 0.0 | 0.231884 | 0.434783 | 0.637681 | 0.942029 | |
| economic_cond_national | 456.0 | 0.559759 | 0.234094 | 0.0 | 0.500000 | 0.500000 | 0.750000 | 1.000000 | |
| economic_cond_household | 456.0 | 0.525219 | 0.228276 | 0.0 | 0.500000 | 0.500000 | 0.750000 | 1.000000 | |
| Blair | 456.0 | 0.571272 | 0.297180 | 0.0 | 0.250000 | 0.750000 | 0.750000 | 1.000000 | |
| Hague | 456.0 | 0.447917 | 0.307687 | 0.0 | 0.250000 | 0.250000 | 0.750000 | 1.000000 | |
| Europe | 456.0 | 0.592982 | 0.338314 | 0.0 | 0.300000 | 0.600000 | 0.925000 | 1.000000 | |
| political_knowledge | 456.0 | 0.534357 | 0.367533 | 0.0 | 0.000000 | 0.666667 | 0.666667 | 1.000000 | |
| gender_male | 456.0 | 0.467105 | 0.499465 | 0.0 | 0.000000 | 0.000000 | 1.000000 | 1.000000 | |

1.4 Apply Logistic Regression and LDA (linear discriminant analysis).

Statistical summary using the statsmodel library

OLS Regression Results

Dep. Variable:

vote

R-squared:

0.391

Model:

OLS

Adj. R-squared:

0.387

Method:

Least Squares

F-statistic:

84.58

Date:

Sun, 10 Dec 2023

Prob (F-statistic):

5.25e-108

Time:

07:32:42

Log-Likelihood:

-402.94

No. Observations:

1061

AIC:

823.9

Df Residuals:

1052

BIC:

868.6

Df Model:

8

Covariance Type: nonrobust

coef

std err

t

P>|t|

[0.025

0.975]

const

0.8157

0.058

13.992

0.000

0.701

0.930

age

-0.1733

0.049

-3.549

0.000

-0.269

-0.078

economic_cond_national

0.3034

0.056

5.423

0.000

0.194

0.413

economic_cond_household

0.0251

0.050

0.507

0.612

-0.072

0.122

Blair

0.3723

0.041

9.047

0.000

0.292

0.453

Hague

-0.4647

0.038

-12.366

0.000

-0.538

-0.391

Europe

-0.2804

0.036

-7.694

0.000

-0.352

-0.209

political_knowledge

-0.1619

0.031

-5.193

0.000

-0.223

-0.101

gender_male

0.0187

0.022

0.839

0.402

-0.025

0.062

Omnibus:

24.713

Durbin-Watson:

2.027

Prob(Omnibus):

0.000

Jarque-Bera (JB):

25.961

Skew:

-0.375

Prob(JB):

2.31e-06

Kurtosis:

2.844

Cond. No.

11.8

Inference:

- The null hypothesis will be rejected for each of the predictor variables and all the predictors will be part of the linear equation as we know that the p value should be less than the significance value of 5% and in the above ols statsmodel we can see that all p- values are much lesser than 5% so we can say that the null hypothesis will be rejected of each predictor variable and all he predictors will be the part of linear equation.
- We need to check the VIF to check the multi -collinearity as if there is multi-collinearity present that it will result in less accuracy and interpretation of coefficients is also being wrong for that we need to must treat the multicollinearity if found. Any VIF score between 1 to 5 is

good to proceed but if it is more than 5 we need to drop that variable so that it will not interrupt in predictions.

VIF of all predicted variables:

```
age                1.015654
economic_cond_national 1.208367
economic_cond_household 1.133530
Blair              1.214217
Hague             1.127716
Europe            1.185442
political_knowledge 1.049427
gender_male       1.039776
dtype: float64
```

We can see that the VIF score is less than 5 so we can say that there is no multi-collinearity in the predicted variables, so we are good to proceed further.

RMSE

Please find below the accuracy metrics/score:

Root mean square error for the testing set is 0.36.

```
The Root Mean Square Error (RMSE) of the model is for testing set is 0.3622858652340443
```

Equation linear equation that depicts individual coefficients of independent variables along with intercept value.

```
usr = 0.8156731537864652 + -0.17333835860235264 * ( age ) + 0.3033682104790156 * ( economic_cond_national ) + 0.025109677108676003 * ( economic_cond_household ) + 0.3723145377382221 * ( Blair ) + -0.4647080673190488 * ( Hague )
```

```
+ -0.2803540095167645 * ( Europe ) + -0.1618600153014007 * ( political_knowledge ) + 0.018690939212194044 * ( gender_male )
```

Classification report of training data and test data.


➞ Classification Report of the training data:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.74 | 0.65 | 0.69 | 307 |
| 1.0 | 0.86 | 0.91 | 0.89 | 754 |
| accuracy | | | 0.83 | 1061 |
| macro avg | 0.80 | 0.78 | 0.79 | 1061 |
| weighted avg | 0.83 | 0.83 | 0.83 | 1061 |

Classification Report of the test data:


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.77 | 0.73 | 0.74 | 153 |
| 1.0 | 0.86 | 0.89 | 0.88 | 303 |
| accuracy | | | 0.83 | 456 |
| macro avg | 0.82 | 0.81 | 0.81 | 456 |
| weighted avg | 0.83 | 0.83 | 0.83 | 456 |

Confusion matrix of training and test data:



AUC for training and test data is shown below along with ROC Curve

➡ AUC for the Training Data: 0.889
AUC for the Test Data: 0.888



Inference:

- Training and Testing results shows that the model is excellent with good precision and recall
- The model is not overfitting or underfitting.

Logistic Regression:

- Split the X and Y into 70:30 training and test data
- Getting the Predicted Classes and Probs

| | 0 | 1 |
|---|----------|----------|
| 0 | 0.199510 | 0.800490 |
| 1 | 0.610744 | 0.389256 |
| 2 | 0.083468 | 0.916532 |
| 3 | 0.030816 | 0.969184 |
| 4 | 0.135993 | 0.864007 |

- Accuracy of training data

0.8303655107778819

- Accuracy of test data

0.8552631578947368

- Classification Report of test and train data


Train data:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.75 | 0.66 | 0.70 | 323 |
| 1.0 | 0.86 | 0.91 | 0.88 | 744 |
| accuracy | | | 0.83 | 1067 |
| macro avg | 0.81 | 0.78 | 0.79 | 1067 |
| weighted avg | 0.83 | 0.83 | 0.83 | 1067 |

- Test data:

| ⊗ | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.80 | 0.68 | 0.73 | 139 |
| 1.0 | 0.87 | 0.92 | 0.90 | 319 |
| accuracy | | | 0.85 | 458 |
| macro avg | 0.83 | 0.80 | 0.81 | 458 |
| weighted avg | 0.85 | 0.85 | 0.85 | 458 |

- AUC for training and test data is shown below along with ROC Curve




- Apply GRID search CV for logistic Regression

Getting probabilities of test set

| | 0 | 1 |
|---|----------|----------|
| 0 | 0.225953 | 0.774047 |
| 1 | 0.090054 | 0.909946 |
| 2 | 0.060530 | 0.939470 |
| 3 | 0.263878 | 0.736122 |
| 4 | 0.018824 | 0.981176 |

- Confusion matrix of test and train data



- Precision of 1's

0.8676470588235294

- Recall of 1's

0.9247648902821317

Inference:

- Training and Testing results shows that the model is excellent with good precision and recall
- The LDA model is better than Logistic regression with better Test accuracy and recall values

Vote refers to original classification in the dataset and predicted vote refers to model output. Age column refers to count grouped by vote and predicted vote to give a comparison. It enables us to see the real mapping of classification across original and prediction in terms of accuracy by each model.

1.5 Apply KNN Model and Naïve Bayes Model. Interpret the results.

KNN Model

- Scaling is a necessary step for KNN Model to proceed . We have already done the scaling in the dataset above and using the scaled data.
- K-Nearest Neighbour is one of the simplest Machine Learning algorithms based on Supervised Learning technique.
- K-NN algorithm assumes the similarity between the new case/data and available cases and put the new case into the category that is most similar to the available categories.
- K-NN algorithm stores all the available data and classifies a new data point based on the similarity. This means when new data appears then it can be easily classified into a well suite category by using K- NN algorithm.
- K-NN algorithm can be used for Regression as well as for Classification but mostly it is used for the Classification problems.
- K-NN is a **non-parametric algorithm**, which means it does not make any assumption on underlying data.

- KNN model is a distance-based mechanism mostly using Euclidean measure. And in this case, we are looking for a classification need and KNN can help building classification based on feature similarity. Since KNN continues to retain the training dataset instead of learning to create a model using the learning, it is a lazy learner and simplest of all models for classification. K in KNN refers to number of nearest neighbours. A small value of K indicates higher influence of noise over result while larger value is cost heavy for compute. So usually a standardized approach is to adopt $K = \sqrt{N}/2$ where N is the size of training data set. Also, K has to be odd number to avoid tie between predicting classes.
- Bayes theorem is an extension of conditional probability that is based on the knowledge of prior probability values of something that has occurred. It provides a mechanism to compute posterior probability of class for the given predictor based on prior probability of the given predictor for the same class. Idea is to factor all available evidence in form of predictors into naïve Bayes rule to obtain more accurate probability for class prediction. Naïve Bayes classifier works on the principle of Bayes theorem with the assumption that input features are independent of each other. Usually not ideal for data sets with large number of numerical attributes, however, does well with noisy and missing data even with low training samples. For those independent variables that are continuous it is assumed to be distributed normally and the estimates go by mean and standard deviation of continuous variable.

We have Scaled data:

| | age | economic_cond_national | economic_cond_household | Blair | Hague | Europe | political_knowledge | gender_male |
|-------|---------------|------------------------|-------------------------|---------------|---------------|---------------|---------------------|---------------|
| count | 1.517000e+03 | 1.517000e+03 | 1.517000e+03 | 1.517000e+03 | 1.517000e+03 | 1.517000e+03 | 1.517000e+03 | 1.517000e+03 |
| mean | -9.367735e-18 | 1.967224e-16 | -1.077290e-16 | 3.747094e-17 | 1.053870e-16 | 4.332578e-17 | -4.215481e-17 | -9.367735e-18 |
| std | 1.000330e+00 | 1.000330e+00 | 1.000330e+00 | 1.000330e+00 | 1.000330e+00 | 1.000330e+00 | 1.000330e+00 | 1.000330e+00 |
| min | -1.926617e+00 | -2.547041e+00 | -2.296796e+00 | -1.988727e+00 | -1.419969e+00 | -1.740556e+00 | -1.421084e+00 | -9.367365e-01 |
| 25% | -8.435773e-01 | -2.781853e-01 | -1.480205e-01 | -1.137217e+00 | -6.083289e-01 | -8.309016e-01 | -1.421084e+00 | -9.367365e-01 |
| 50% | -7.907881e-02 | -2.781853e-01 | -1.480205e-01 | 5.658022e-01 | -6.083289e-01 | -2.244654e-01 | 4.238322e-01 | -9.367365e-01 |
| 75% | 8.128361e-01 | 8.562424e-01 | 9.263674e-01 | 5.658022e-01 | 1.014951e+00 | 9.884072e-01 | 4.238322e-01 | 1.067536e+00 |
| max | 2.469250e+00 | 1.990670e+00 | 2.000755e+00 | 1.417312e+00 | 1.826592e+00 | 1.291625e+00 | 1.346290e+00 | 1.067536e+00 |


- Dataset has been split into 70 :30 ratio (Train and test respectively)

- We have taken 5 as K in the below analysis.
- Confusion matrix and classification report on train data.

```
(X) [[220 107]
      [ 61 749]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.67 | 0.72 | 327 |
| 1 | 0.88 | 0.92 | 0.90 | 810 |
| accuracy | | | 0.85 | 1137 |
| macro avg | 0.83 | 0.80 | 0.81 | 1137 |
| weighted avg | 0.85 | 0.85 | 0.85 | 1137 |

- AUC ROC Curve KNN Train data




- Confusion matrix and classification report on test data.

```
[[ 85  48]
 [ 27 220]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.76 | 0.64 | 0.69 | 133 |
| 1 | 0.82 | 0.89 | 0.85 | 247 |
| accuracy | | | 0.80 | 380 |
| macro avg | 0.79 | 0.76 | 0.77 | 380 |
| weighted avg | 0.80 | 0.80 | 0.80 | 380 |

- AUC ROC Curve KNN Test data

the auc curve 0.848



- For K =7 the below analysis has been done.
- Confusion matrix and classification report on train data.

```
0.8434476693051891
[[212 115]
 [ 63 747]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.77 | 0.65 | 0.70 | 327 |
| 1 | 0.87 | 0.92 | 0.89 | 810 |
| accuracy | | | 0.84 | 1137 |
| macro avg | 0.82 | 0.79 | 0.80 | 1137 |
| weighted avg | 0.84 | 0.84 | 0.84 | 1137 |

- Confusion matrix and classification report on test data.

```
0.8052631578947368
[[ 82  51]
 [ 23 224]]
```


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.62 | 0.69 | 133 |
| 1 | 0.81 | 0.91 | 0.86 | 247 |
| accuracy | | | 0.81 | 380 |
| macro avg | 0.80 | 0.76 | 0.77 | 380 |
| weighted avg | 0.80 | 0.81 | 0.80 | 380 |

- AUC score

```
[0.23947368421052628,
0.21052631578947367,
0.19736842105263153,
0.1947368421052632,
0.18947368421052635,
0.19210526315789478,
0.20789473684210524,
0.19210526315789478,
0.20263157894736838,
0.20263157894736838]
```


- AUC ROC curve after n classifier for train data set

```
➡ the auc curve 0.889
[<matplotlib.lines.Line2D at 0x7c88febf0be0>]
```



- AUC ROC curve after n classifier for test data set

the auc curve 0.885
[<matplotlib.lines.Line2D at 0x7c890153e3e0>]



• Naive Bayes Model

- Taking 70:30 training and test data
- Predicted class is set to be 1 for true and 0 for false
- Random number seeding is set to be 7 for repeatability of code
- Fitting of model has taken place

Confusion matrix and classification report has been shown below for test data

```
GaussianNB()
      precision    recall  f1-score   support

    0.0         0.72     0.71     0.72         147
    1.0         0.86     0.87     0.87         309

 accuracy         0.82         456
 macro avg         0.79     0.79     0.79         456
 weighted avg         0.82     0.82     0.82         456

[[105  42]
 [ 41 268]]
```

Confusion matrix and classification report has been shown below for train data.

```
GaussianNB()
precision    recall  f1-score   support

     0       0.72     0.69     0.71       313
     1       0.87     0.89     0.88       748


 accuracy          0.83       1061
 macro avg         0.80     0.79     0.79       1061
 weighted avg      0.83     0.83     0.83       1061

[[216  97]
 [ 82 666]]
```

- We have built decision tree model
- Scoring of our decision tree

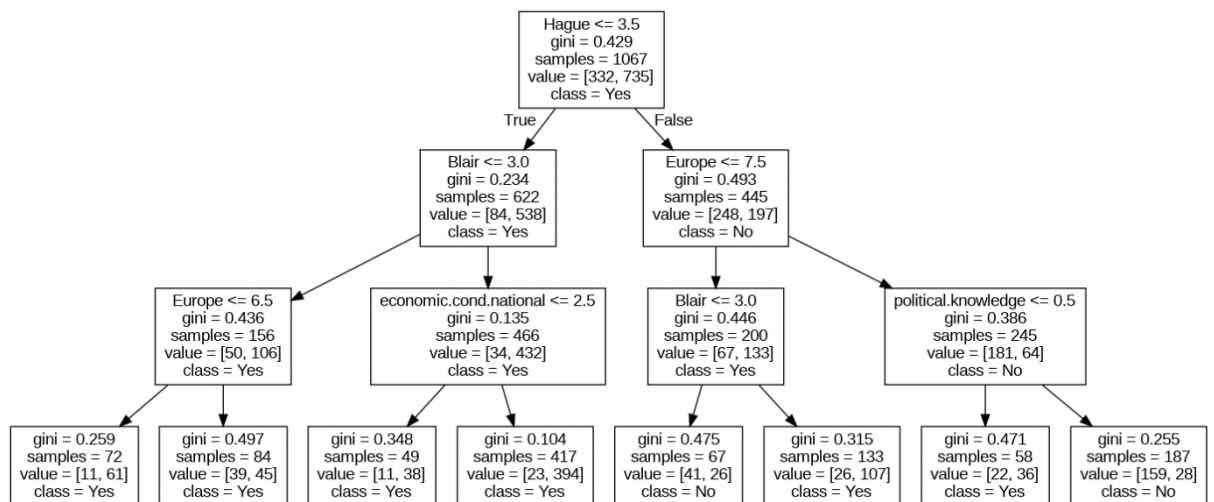
```
1.0
0.7620087336244541
```

- Visualization of the decision tree



- Reducing over fitting (Regularization)

```
0.8256794751640113
0.8056768558951966
```




- The importance of features in the tree building (The importance of a feature is computed as the $\#(\text{normalized})$ total reduction of the criterion brought by that feature. It is also known as the Gini importance)

| | Imp |
|-------------------------|----------|
| age | 0.000000 |
| economic_cond_national | 0.000000 |
| economic_cond_household | 0.016012 |
| Blair | 0.291569 |
| Hague | 0.514417 |
| Europe | 0.139003 |
| political_knowledge | 0.038999 |
| gender_male | 0.000000 |

Confusion matrix for d tree

0.7894736842105263
<Axes: >



Inference:

1. Training and Testing results shows that the model neither overfitting nor underfitting.
2. The Naive Bayes model also performs well with better accuracy and recall values.
3. Even though NB and KNN have same Train and Test accuracy. Based on their recall value in test dataset it is evident that KNN performs better than Naive Bayes.

1.6 Model Tuning, Bagging (Random Forest should be applied for Bagging) and Boosting.

Model tuning means finding better parameters for the model and not just use the default values. For what parameters can be changed refer to the algorithm documentation.

Each ML algorithm has its own parameters to tune the model. Please check the grid search implementation in mentoring session notebooks. Ex: In random forest, two important hyperparameters are `n_estimators`, `max_features`

- Applying GridsearchCV as model tuning.
- Prediction on test data

```
array([0., 1., 1., ..., 1., 1., 1.])
```

- Prediction on test data

```

array([1., 1., 1., 0., 1., 1., 0., 1., 1., 1., 1., 1., 1., 0., 1., 1., 1.,
       1., 0., 1., 1., 0., 0., 1., 1., 1., 1., 1., 1., 0., 1., 0., 0., 1.,
       0., 1., 0., 1., 0., 1., 1., 1., 1., 0., 1., 0., 1., 1., 0., 1., 1.,
       1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 1., 1., 1., 1., 1.,
       0., 1., 0., 0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0.,
       1., 1., 1., 0., 1., 0., 0., 1., 1., 0., 1., 1., 0., 1., 1., 0., 1.,
       1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0., 0., 1., 0., 1., 1.,
       0., 1., 1., 0., 0., 1., 1., 1., 0., 1., 1., 1., 1., 1., 0., 1.,
       0., 1., 0., 1., 1., 0., 1., 1., 1., 1., 0., 0., 1., 0., 0., 1., 1.,
       1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 1., 0., 1., 0., 1., 1., 0.,
       0., 0., 1., 1., 1., 0., 1., 0., 1., 1., 0., 1., 0., 1., 1., 0.,
       1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 0., 1., 1., 0., 1., 1.,
       0., 1., 1., 1., 1., 0., 1., 1., 1., 1., 0., 1., 0., 1., 1., 0.,
       1., 1., 1., 1., 1., 0., 0., 0., 1., 1., 1., 1., 0., 1., 1., 1.,
       1., 0., 1., 1., 1., 1., 0., 0., 1., 1., 0., 1., 1., 1., 1.,
       0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 1., 1.,
       0., 1., 0., 1., 1., 1., 1., 0., 1., 1., 1., 0., 0., 1., 1., 0.,
       1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 1., 1.,
       0., 1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
       0., 1., 0., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1.])

```

Accuracy on train

0.8727615457115928

Accuracy on test data

0.8377192982456141

Generating confusion matrix of test and train respectively

```

array([[211, 96],
       [ 39, 715]])

```

```

array([[101, 52],
       [ 22, 281]])

```

Generate a classification report for the training and test data

Training data

| | | | | | | | | | | | | | | | | | | | |
|----|-----|-----------|--------|----------|-------------|--------------|------|------|------|--------|-----|------|------|------|---------|----------|------|--------|-----|
| | | precision | recall | f1-score | support\n\n | 0.0 | 0.84 | 0.69 | 0.76 | 307\n | 1.0 | 0.88 | 0.95 | 0.91 | 754\n\n | accuracy | 0.87 | 1061\n | mac |
| ro | avg | 0.86 | 0.82 | 0.84 | 1061\n | weighted avg | 0.87 | 0.87 | 0.87 | 1061\n | | | | | | | | | |


Test data

| | | | | | | | | | | | | | | | | | | | |
|----|-----|-----------|--------|----------|-------------|--------------|------|------|------|-------|-----|------|------|------|---------|----------|------|-------|-----|
| | | precision | recall | f1-score | support\n\n | 0.0 | 0.82 | 0.66 | 0.73 | 153\n | 1.0 | 0.84 | 0.93 | 0.88 | 303\n\n | accuracy | 0.84 | 456\n | mac |
| ro | avg | 0.83 | 0.79 | 0.81 | 456\n | weighted avg | 0.84 | 0.84 | 0.83 | 456\n | | | | | | | | | |

Bagging:

Bagging, also known as bootstrap aggregation, is the ensemble learning method that is commonly used to reduce variance within a noisy dataset. In bagging, a random sample of data in a training set is selected with replacement—meaning that the individual data points can be chosen more than once.

- N_estimator is taken to be 50 and random state as 1
- Confusion matrix for bagging



- Bagging model score , Confusion matrix and classification report for train

```

1.0
[[307  0]
 [ 0 754]]

```


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 307 |
| 1 | 1.00 | 1.00 | 1.00 | 754 |
| accuracy | | | 1.00 | 1061 |
| macro avg | 1.00 | 1.00 | 1.00 | 1061 |
| weighted avg | 1.00 | 1.00 | 1.00 | 1061 |

AUC ROC Curve for train


```

AUC: 1.000
[<matplotlib.lines.Line2D at 0x7d08078062c0>]

```



AUC ROC Curve for test



- Bagging model score , Confusion matrix and classification report for test

```
0.8201754385964912
[[108  45]
 [ 37 266]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.74 | 0.71 | 0.72 | 153 |
| 1 | 0.86 | 0.88 | 0.87 | 303 |
| accuracy | | | 0.82 | 456 |
| macro avg | 0.80 | 0.79 | 0.80 | 456 |
| weighted avg | 0.82 | 0.82 | 0.82 | 456 |

- Boosting – ADA Boosting

Confusion matrix of train data

```
array([[238,  94],  
       [ 69, 666]], dtype=int64)
```

Confusion matrix of test data

```
array([[ 90,  40],  
       [ 43, 285]], dtype=int64)
```

accuracy score for the training and test data

train

```
0.8472352389878163
```

Test

```
0.8187772925764192
```

- ADA Bosting score , Confusion matrix and classification report for train

0.8472352389878163

[[238 94]


[69 666]]

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.72 | 0.74 | 332 |
| 1 | 0.88 | 0.91 | 0.89 | 735 |
| accuracy | | | 0.85 | 1067 |
| macro avg | 0.83 | 0.81 | 0.82 | 1067 |
| weighted avg | 0.84 | 0.85 | 0.85 | 1067 |


- AUC-ROC of train data

AUC: 0.915

[<matplotlib.lines.Line2D at 0x7d0804e561a0>]




- AUC-ROC of test data



- Gradient Boosting

Gradient boosting is a machine learning technique used in regression and classification tasks, among others. It gives a prediction model in the form of an ensemble of weak prediction models, i.e., models that make very few assumptions about the data, which are typically simple decision trees.


Confusion matrix of test:



Random forest

A random forest (RF) is an ensemble of decision trees in which each decision tree is trained with a specific random noise. Random forests are the most popular form of decision tree ensemble.

Confusion matrix of test data:



Inference:

1. Bagging model performs the best with 95% train accuracy. And also have 96% precision and above 90% recall which is better than any other models that we have performed in here with the Election dataset.
2. Rest all the models are more or less have same accuracy of 89%

1.7 Performance Metrics: Check the performance of Predictions on Train and Test sets using Accuracy, Confusion Matrix, Plot ROC curve and get ROC_AUC score for each model. Final Model: Compare the models and write inference which model is best/optimized.

LDA

Classification report of training data and test data.


→ Classification Report of the training data:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.74 | 0.65 | 0.69 | 307 |
| 1.0 | 0.86 | 0.91 | 0.89 | 754 |
| accuracy | | | 0.83 | 1061 |
| macro avg | 0.80 | 0.78 | 0.79 | 1061 |
| weighted avg | 0.83 | 0.83 | 0.83 | 1061 |

Classification Report of the test data:


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.77 | 0.73 | 0.74 | 153 |
| 1.0 | 0.86 | 0.89 | 0.88 | 303 |
| accuracy | | | 0.83 | 456 |
| macro avg | 0.82 | 0.81 | 0.81 | 456 |
| weighted avg | 0.83 | 0.83 | 0.83 | 456 |

Confusion matrix of training and test data:



AUC for training and test data is shown below along with ROC Curve

→ AUC for the Training Data: 0.889
AUC for the Test Data: 0.888



Logistic Regression:


Train data:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.75 | 0.66 | 0.70 | 323 |
| 1.0 | 0.86 | 0.91 | 0.88 | 744 |
| accuracy | | | 0.83 | 1067 |
| macro avg | 0.81 | 0.78 | 0.79 | 1067 |
| weighted avg | 0.83 | 0.83 | 0.83 | 1067 |


- Test data:

| ⊗ | | precision | recall | f1-score | support |
|---|--------------|-----------|--------|----------|---------|
| | 0.0 | 0.80 | 0.68 | 0.73 | 139 |
| | 1.0 | 0.87 | 0.92 | 0.90 | 319 |
| | accuracy | | | 0.85 | 458 |
| | macro avg | 0.83 | 0.80 | 0.81 | 458 |
| | weighted avg | 0.85 | 0.85 | 0.85 | 458 |

- AUC for training and test data is shown below along with ROC Curve



- Apply GRID search CV for logistic Regression
- Confusion matrix of test and train data



KNN Model

- Confusion matrix and classification report on train data.

```

[[220 107]
 [ 61 749]]
precision    recall  f1-score   support


0           0.78     0.67     0.72       327
1           0.88     0.92     0.90       810

accuracy          0.85       1137
macro avg         0.83     0.80     0.81       1137
weighted avg      0.85     0.85     0.85       1137

```

- AUC ROC Curve KNN Train data

the auc 0.924




- Confusion matrix and classification report on test data.

```
[[ 85  48]
 [ 27 220]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.76 | 0.64 | 0.69 | 133 |
| 1 | 0.82 | 0.89 | 0.85 | 247 |
| accuracy | | | 0.80 | 380 |
| macro avg | 0.79 | 0.76 | 0.77 | 380 |
| weighted avg | 0.80 | 0.80 | 0.80 | 380 |

- AUC ROC Curve KNN Test data

the auc curve 0.848



- For K =7 the below analysis has been done.
- Confusion matrix and classification report on train data.

```
0.8434476693051891
[[212 115]
 [ 63 747]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.77 | 0.65 | 0.70 | 327 |
| 1 | 0.87 | 0.92 | 0.89 | 810 |
| accuracy | | | 0.84 | 1137 |
| macro avg | 0.82 | 0.79 | 0.80 | 1137 |
| weighted avg | 0.84 | 0.84 | 0.84 | 1137 |


- Confusion matrix and classification report on test data.

```
0.8052631578947368
[[ 82  51]
 [ 23 224]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.62 | 0.69 | 133 |
| 1 | 0.81 | 0.91 | 0.86 | 247 |
| accuracy | | | 0.81 | 380 |
| macro avg | 0.80 | 0.76 | 0.77 | 380 |
| weighted avg | 0.80 | 0.81 | 0.80 | 380 |


- AUC ROC curve after n classifier for train data set

→ the auc curve 0.889
[<matplotlib.lines.Line2D at 0x7c88feb0be0>]



- AUC ROC curve after n classifier for test data set

→ the auc curve 0.885
[<matplotlib.lines.Line2D at 0x7c890153e3e0>]



- **Naive Bayes Model**

Confusion matrix and classification report has been shown below for test data

```

GaussianNB()
              precision    recall  f1-score   support

    0.0       0.72       0.71       0.72         147
    1.0       0.86       0.87       0.87         309

 accuracy          0.82         456
 macro avg         0.79         456
 weighted avg      0.82         456

[[105  42]
 [ 41 268]]

```

Confusion matrix and classification report has been shown below for train data.

```

GaussianNB()
              precision    recall  f1-score   support

    0       0.72       0.69       0.71         313
    1       0.87       0.89       0.88         748

 accuracy          0.83        1061
 macro avg         0.80         1061
 weighted avg      0.83         1061

[[216  97]
 [ 82 666]]

```

Confusion matrix for d tree

```

0.7894736842105263
<Axes: >

```



- ```
array([0., 1., 1., ..., 1., 1., 1.])
```

- [illegible]

0.8727615457115928

0.8377192982456141



```
array([[211, 96],
 [39, 715]])
```


---

```
array([[101, 52],
 [22, 281]])
```

```
array([[101, 52],
 [22, 281]])
```

- Confusion matrix for bagging

- Confusion matrix for bagging




- Bagging model score , Confusion matrix and classification report for train

```
1.0
[[307 0]
 [0 754]]
```


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 307     |
| 1            | 1.00      | 1.00   | 1.00     | 754     |
| accuracy     |           |        | 1.00     | 1061    |
| macro avg    | 1.00      | 1.00   | 1.00     | 1061    |
| weighted avg | 1.00      | 1.00   | 1.00     | 1061    |

## AUC ROC Curve for train

⇒ AUC: 1.000  
[<matplotlib.lines.Line2D at 0x7d08078062c0>]



## AUC ROC Curve for test



- Bagging model score , Confusion matrix and classification report for test

```
0.8201754385964912
[[108 45]
 [37 266]]
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.74      | 0.71   | 0.72     | 153     |
| 1            | 0.86      | 0.88   | 0.87     | 303     |
| accuracy     |           |        | 0.82     | 456     |
| macro avg    | 0.80      | 0.79   | 0.80     | 456     |
| weighted avg | 0.82      | 0.82   | 0.82     | 456     |

- Boosting – ADA Boosting

Confusion matrix of train data

```
array([[238, 94],
 [69, 666]], dtype=int64)
```

Confusion matrix of test data

```
array([[90, 40],
 [43, 285]], dtype=int64)
```

accuracy score for the training and test data

train

```
0.8472352389878163
```

Test

```
0.8187772925764192
```

- ADA Bosting score , Confusion matrix and classification report for train

```
0.8472352389878163
```

```
[[238 94]
```


```
 [69 666]]
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.78      | 0.72   | 0.74     | 332     |
| 1            | 0.88      | 0.91   | 0.89     | 735     |
| accuracy     |           |        | 0.85     | 1067    |
| macro avg    | 0.83      | 0.81   | 0.82     | 1067    |
| weighted avg | 0.84      | 0.85   | 0.85     | 1067    |


- AUC-ROC of train data

```
AUC: 0.915
```

```
[<matplotlib.lines.Line2D at 0x7d0804e561a0>]
```




- AUC-ROC of test data



- Gradient Boosting

Gradient boosting is a machine learning technique used in regression and classification tasks, among others. It gives a prediction model in the form of an ensemble of weak prediction models, i.e., models that make very few assumptions about the data, which are typically simple decision trees.


Confusion matrix of test:



## Random forest

A random forest (RF) is an ensemble of decision trees in which each decision tree is trained with a specific random noise. Random forests are the most popular form of decision tree ensemble.

Confusion matrix of test data:



## 1.8 Based on these predictions, what are the insights?

1)Dataset has 10 variables. Unnamed: 0 is serial number so we can drop that.

2)We see that variable vote and gender contain string value.

3)We know that modelling cannot take string value. To perform the analysis, we will be converting string value to integer data type.

4)Age refers to the person who have entered/ gained voting rights. The maximum economic condition of the nation falls under the moderate category. The maximum economic condition of the household goods falls under the moderate category. Blair is the labour leader.

5)Model score for all the models seem to be similar for both the training and the test sets. AUC scores are also same for almost all the models. From the confusion matrix, we see that actual and predicted data are very close to each other. This is the sign of right fit model. We see even the F1 scores are almost same on all the models. Bagging and boosting gave out excellent results proving there is no overfitting or underfitting. We see that Tuning the model has optimized the results for the training set and hence we can consider this for ADA Boost. As per the models, and the best-case scenario given accuracy of more than 83%, we can say that in between 57 to 59% of seats are reserved for Labour party.

In the worst-case scenario, given accuracy of more than 82%, a minimum of in between 56 to 58% of seats are reserved for the Conservative Party.

Considering other variables, in the model, we can say that there are multiple fence sitters who can vote for either of the parties by getting influenced from various variables in action. Out of voters choosing labour party, more than 50% are female and less than 49% are male. Labour party can see to increase influence on males to increase its vote bank.



## Problem 2:

In this particular project, we are going to work on the inaugural corpora from the nltk in Python. We will be looking at the following speeches of the Presidents of the United States of America:

1. President Franklin D. Roosevelt in 1941
2. President John F. Kennedy in 1961
3. President Richard Nixon in 1973

(Hint: use `.words()`, `.raw()`, `.sent()` for extracting counts)

### 2.1 Find the number of characters, words, and sentences for the mentioned documents.

The number of character, words and sentence in the below documents

- Franklin D. Roosevelt
- John F. Kennedy
- Richard Nixon

| Speech                | Chars | Words | Sentences |
|-----------------------|-------|-------|-----------|
| Franklin D. Roosevelt | 7571  | 1536  | 68        |
| John F. Kennedy       | 7618  | 1546  | 52        |
| Richard Nixon         | 9991  | 2028  | 69        |

### 2.2) Remove all the stop words from the three speeches. Show the word count before and after the removal of stop words. Show a sample sentence after the removal of stop words.

Before the removal of stop words:

| Speech                | Words |
|-----------------------|-------|
| Franklin D. Roosevelt | 1536  |
| John F. Kennedy       | 1546  |
| Richard Nixon         | 2028  |

- **After the removal of stopwords**

Number of words in Speech of President Franklin D. Roosevelt in 1941-after removal of stopwords are 632  
Number of words in Speech of President John F. Kennedy in 1961-after removal of stopwords are 697  
Number of words in Speech of President Richard Nixon in 1973-after removal of stopwords are 836


2.3) Which word occurs the most number of times in his inaugural address for each president? Mention the top three words. (after removing the stopwords)

```
Top word which occurred the most number of times in inaugural address of President Franklin D. Roosevelt in 1941 is [('nation', 12)]
Top word which occurred the most number of times in inaugural address of President John F. Kennedy in 1961 is [('let', 16)]
Top word which occurred the most number of times in inaugural address of President Richard Nixon in 1973 is [('us', 26)]
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
```


```
Top 3 words which occurred the most number of times in inaugural address of President Franklin D. Roosevelt in 1941 is [('nation', 12), ('know', 10), ('spirit', 9)]
Top 3 words which occurred the most number of times in inaugural address of President John F. Kennedy in 1961 is [('let', 16), ('us', 12), ('world', 8)]
Top 3 words which occurred the most number of times in inaugural address of President Richard Nixon in 1973 is [('us', 26), ('let', 22), ('america', 21)]
```

2.4 Plot the word cloud of each of the speeches of the variable. (after removing the stopwords)

Word Cloud for President Franklin D. Roosevelt speech in 1941



## Word Cloud for President John F. Kennedy speech in 1961



### Word Cloud for President Richard Nixon speech in 1973

